ChatGPT

Key Concepts

- **Interval Forecasting vs Point Forecasting:** Point forecasts provide single-valued predictions (often the mean or median) but ignore uncertainty [1]. Interval forecasting yields a range (prediction interval) within which the future value is expected to lie with a given probability (e.g. 95%) [2]. By capturing uncertainty, interval forecasts convey the confidence or risk associated with a prediction, which is crucial in volatile markets. In practice, models tend to underestimate uncertainty (intervals often too narrow) if they don't account for all sources of variability [3]. Interval forecasts (and more generally **probabilistic forecasts**) thus provide a fuller picture by estimating the entire distribution of future outcomes, enabling risk assessment beyond a single-point estimate [4]. For example, a 72-hour **prediction interval** for a token return might indicate a high likelihood of extreme moves, guiding traders to adjust positions accordingly [5] [6].

- **Quantile Regression (Linear, Boosted, Kernel):** Quantile regression (QR) generalizes linear regression to model specified quantiles (percentiles) of the response distribution, instead of the mean [7]. The seminal work of Koenker and Bassett (1978) introduced linear QR, which estimates conditional quantiles by minimizing "pinball" loss (an asymmetric absolute loss) [7]. This allows different slope coefficients for, say, the median vs. the 95th-percentile response, capturing heteroskedasticity or skewed residuals better than mean regression. **Boosted quantile regression** uses ensemble techniques (e.g. gradient boosting) with quantile loss to fit complex, non-linear quantile functions. Tree-based gradient boosting (as in XGBoost or LightGBM) can directly optimize the pinball loss for chosen quantiles, yielding fast, non-parametric quantile estimates. **Kernel quantile regression** (KQR) takes a non-parametric approach using kernel smoothing in feature space [8]. For instance, kernel QR (e.g. using Gaussian kernels in a reproducing kernel Hilbert space) fits quantiles without a fixed parametric form, which can capture nonlinear relationships in financial time series. All these quantile regression methods share a key benefit: they produce **prediction intervals** and **VaR estimates** directly (since a prediction interval is essentially two estimated quantiles, e.g. 2.5% and 97.5%) [5], making them valuable for risk-aware forecasting.

- **Quantile Regression Forests (QRF):** QRF is an extension of the random forest ensemble that predicts the full conditional distribution of a response rather than just the mean [9]. In a standard random forest, averaging tree outputs gives an accurate mean prediction; QRF instead **retains the distribution of target values** in the leaf nodes across all trees [10]. By collecting the set of all training observations that fall into the same leaves as a query point, QRF can estimate any conditional quantile (e.g. 5th, 50th, 95th percentile) of the response [11]. This non-parametric approach is well-suited for high-dimensional financial data, as it makes minimal assumptions about return distributions [12]. For example, QRF can output a 95% interval of a 3-day return forecast by providing the 2.5th and 97.5th percentile predictions, inherently capturing volatility and skewness. Originally proposed by Meinshausen (2006) [9], QRF has been shown to be **consistent** and often competitive in predictive power [13]. Its ability to model **tail behavior** without assuming a Gaussian distribution is a major advantage in crypto markets, where returns are fat-tailed. Recent studies leverage QRF and related **generalized random forests** to forecast Value-at-Risk of crypto assets, finding them superior to traditional parametric models under high volatility [14] [15].

- **Tail-Risk Metrics & Risk-Adjusted Sizing:** Tail-risk metrics quantify the risk of extreme outcomes (usually losses) in an asset's return distribution. A classic metric is **Value-at-Risk (VaR)** – the loss threshold that will not be exceeded at a given confidence level. For example, the 10-day 99% VaR is the loss level which has only a 1% chance of being exceeded in 10 days [16]. Another is **Expected Shortfall (ES)** (or Conditional VaR), which asks "if things do get bad, what is our expected loss?" – i.e. the average loss in the worst q% cases [17]. These measures are directly related to quantiles: VaR is essentially a low-order quantile of the loss distribution, and ES is a tail-average beyond that quantile. **Risk-adjusted position sizing** means dynamically scaling trade sizes based on risk estimates so that portfolio risk remains consistent. For instance, one can target a constant VaR by reducing position size when forecasted volatility or VaR is high, and vice versa. Recent techniques propose using **volatility estimates** or **EVT-based CVaR** to set trade sizes such that a fixed tail-risk level is maintained [18] [19]. One algorithm adjusts position sizes using Extreme Value Theory (EVT) to estimate CVaR, ensuring the probability of large drawdowns stays under control [20]. In practice, a trader might allocate capital such that the 95% VaR of a 72-hour Solana position equals some fraction of the portfolio. By doing so, the **tail risk** (e.g. potential 5% worst-case loss) is kept constant regardless of market volatility. This approach aligns position sizing with forecasted risk, leading to more **tail-risk-aware** strategies than traditional fixed-size trades.

- **Multi-Day Swing-Trade Horizon Considerations:** Forecasting over a multi-day "swing trade" horizon (e.g. 3–5 days) introduces different challenges than daily or intraday forecasting. Longer horizons tend to accumulate more uncertainty – forecast error generally **increases as the horizon extends** [21]. This is partly due to compounding volatility and the possibility of regime shifts over several days. In a 72-hour horizon, one must account for the fact that crypto markets trade 24/7; a three-day return might span a weekend or specific market events, affecting its distribution. Multi-day forecasts can be obtained via **iterated one-day models** (compounding one-day predictions) or **direct multi-step models**. Direct models (predicting the k-day return directly) may better capture multi-day trends or mean-reversion that one-day models miss, but they face data limitations since overlapping multi-day returns reduce effective sample size. Swing trading aims to capture intermediate trends, so **forecasting swing returns often involves trade-offs**: short-term noise is smoothed out, but one must predict over potentially varying market conditions. The horizon also affects calibration of intervals – a 3-day 95% interval will typically be wider than a 1-day interval, often scaling with $\sqrt{3}$ times volatility under a diffusion assumption (though in practice heavy tails and autocorrelation can make the scaling non-linear). Additionally, position management differs: a swing trader might hold a position over multiple days, so using interval forecasts, they might decide to **tighten stop-losses or reduce leverage** if the 3-day 5% VaR is above a threshold. Overall, multi-day forecasting requires methods that handle **path dependency** and **compounding risk**, and evaluations should consider metrics like coverage of multi-day intervals or profitability of using multi-day forecast bands in swing trading strategies.

Seminal & Recent Publications by Category

(a) Non-Parametric Interval Forecasting in Equities/FX: *Pioneering and recent works using distribution-free or semi-parametric methods to predict intervals or densities in traditional markets.*

- **Tay, A.S. & Wallis, K.F. (2000). "Density Forecasting: A Survey."** ***Journal of Forecasting*** **19(4): 235–254. DOI: 10.1002/1099-131X(200007)19:4<235::AID-FOR754>3.0.CO;2-R.**
 Summary: A comprehensive survey of density forecasting techniques [5] [6]. It highlights the importance of going beyond point forecasts to estimate the entire predictive distribution. The paper reviews methods including historical simulation, non-parametric kernel density forecasts, and the evaluation of interval forecasts (coverage tests), laying groundwork for later VaR and interval prediction research in finance.

- **Engle, R.F. & Manganelli, S. (2004). "CAViaR: Conditional Autoregressive Value at Risk by Regression Quantiles."** ***Journal of Business & Economic Statistics*** **22(4): 367–381. DOI: 10.1198/073500104000000370.**
 Summary: Introduces the CAViaR model, which directly models the quantile of asset returns (VaR) as an autoregressive process [22]. Rather than assuming a parametric distribution for returns, CAViaR uses quantile regression to update VaR estimates over time. This semi-parametric approach was seminal in showing that **non-parametric quantile methods** can outperform GARCH for VaR forecasting, and it paved the way for many subsequent non-parametric risk models.

- **Chen, S.X. & Tang, C.Y. (2005). "Nonparametric Inference of Value-at-Risk for Dependent Financial Returns."** ***Journal of Financial Econometrics*** **3(2): 227–255. DOI: 10.1093/jjfinec/nbi009.**
 Summary: This paper develops a non-parametric kernel estimator for VaR under serial dependence. The authors use a double-kernel local linear quantile estimator to account for lagged returns and conditional heteroskedasticity [23]. Applied to stock index returns, their method produces competitive VaR forecasts without assuming a specific distribution. It is an early example of using **kernel quantile regression** in finance, demonstrating improved accuracy in capturing tail risk over moving-window historical simulation.

- **Schaumburg, J. (2012). "Predicting Extreme Value-at-Risk: Nonparametric Quantile Regression with EVT Refinements." (Discussion Paper, Humboldt U.).**
 Summary: Proposes a hybrid of non-parametric quantile regression and Extreme Value Theory (EVT) for predicting very extreme VaR (e.g. 0.1% tail) [24] [25]. Moderate tails (1–5% VaR) are estimated via kernel quantile regression, while the most extreme tail is modeled by a Generalized Pareto distribution fit to residuals. In equity index data, this approach yielded more accurate tail risk forecasts than pure GARCH or HS (historical simulation), underscoring the value of combining **distribution-free quantile estimates** with EVT for rare-event risk.

- **Christoffersen, P. (1998). "Evaluating Interval Forecasts."** ***International Economic Review*** **39(4): 841–862. DOI: 10.2307/2527343.**
 Summary: While not a forecasting method per se, this highly cited paper set the framework for backtesting prediction intervals in finance. Christoffersen introduced tests for correct interval coverage and independence of hits (exceptions), which became standard for evaluating VaR models. By formalizing how to judge **interval forecast accuracy** (e.g.

Kupiec's test for coverage), this work indirectly spurred improvements in interval forecasting methods (like non-parametric approaches) to meet rigorous evaluation standards.

(b) Quantile Regression Forests (QRF) and ML Quantile Methods: *Key papers on ensemble and machine learning approaches for conditional quantile estimation.*

- **Meinshausen, N. (2006). "Quantile Regression Forests."** ***Journal of Machine Learning Research*** **7(Jun): 983–999.**
 Summary: The foundational paper introducing Quantile Regression Forests [9]. It proved that random forests can recover the full conditional distribution of $Y|X$ by keeping the distribution of responses in each leaf. Meinshausen's QRF algorithm estimates conditional quantiles with no parametric assumptions and is shown to be consistent. This paper includes examples where QRF provides **accurate prediction intervals** that adapt to heteroskedasticity (e.g. varying interval widths when data dispersion varies with X) [26] [27]. It has since been widely applied to probabilistic forecasting problems, including financial risk predictions.

- **Athey, S., Tibshirani, J., & Wager, S. (2019). "Generalized Random Forests."** ***Annals of Statistics*** **47(2): 1148–1178. DOI: 10.1214/18-AOS1709.**
 Summary: Introduces a broad framework that includes quantile regression forests as a special case. Generalized Random Forests (GRF) use forest weightings for estimation of general quantities (treatment effects, quantiles, etc.). The paper provides theoretical foundations (consistency, asymptotic normality) for forest-based estimators of conditional quantiles. This has influenced financial applications by allowing, for instance, **quantile treatment effect** estimation and more robust **causal inference** in economics, and it underpins advanced VaR forecasting approaches (e.g. GRF for crypto VaR [28]).

- **Görgen, K., Buse, R., & Schienle, M. (2022). "Predicting Value-at-Risk for Cryptocurrencies with Generalized Random Forests."** ***arXiv preprint*** **arXiv:2203.08224.**
 Summary: A recent study that applies ML quantile methods to crypto risk. The authors use a quantile version of GRF to forecast daily VaR of 105 cryptocurrencies [14]. They find that the GRF-based approach **outperforms classical quantile models** (like GARCH and CAViaR) in out-of-sample VaR prediction, especially during turbulent periods [29]. This paper demonstrates the efficacy of tree-based ensemble quantile methods in capturing the heavy-tailed, non-linear dynamics of crypto returns, validating QRF/GRF as cutting-edge tools for risk forecasting.

- **Xie, Y. & Shi, H. (2019). "Time Series Quantile Regression with Random Forests."** ***Journal of Statistical Theory and Practice*** **13, 18. DOI: 10.1007/s42519-018-0002-1.**
 Summary: Investigates the use of random forests for quantile regression in a time-series context. The authors modify QRF to handle lagged observations as features and apply it to macroeconomic and financial time series. They show improvement over linear quantile AR models in terms of interval forecast accuracy. This work is notable for addressing the specific challenges of **dependence and autocorrelation** in time series quantile prediction, bridging a gap between ML quantile methods and traditional time-series econometrics.

- **Bauer, I., Haupt, H., & Linner, S. (2024). "Pinball Boosting of Regression Quantiles."** ***Computational Statistics & Data Analysis*** **200: 108027. DOI: 10.1016/j.csda.2024.108027.**
 Summary: Presents a boosting algorithm explicitly tailored for quantile regression (often nicknamed "pinball boosting" after the quantile loss function). Through simulation and empirical tests, the authors show that boosted ensembles of shallow trees can estimate extreme quantiles more accurately than single-model methods. This is relevant for finance as it provides a way to build **additive quantile models** that capture complex interactions. The paper finds pinball boosting competitive with classical methods, highlighting how boosting can improve predictive **sharpness** (narrower, well-calibrated intervals) in forecasting applications.

(c) Linear Quantile Regression & Boosting-Based Approaches: *Notable works on classical quantile regression, its extensions, and related boosting methods.*

- **Koenker, R. & Bassett, G. (1978). "Regression Quantiles."** ***Econometrica*** **46(1): 33–50. DOI: 10.2307/1913643.**
 Summary: The original paper that introduced quantile regression. It defines the regression quantile estimator and shows that by minimizing asymmetrically weighted absolute errors one can estimate any conditional quantile of $Y|X$. This laid the theoretical groundwork for all linear quantile regression applications in economics and finance. Its influence is vast – e.g., in finance it enabled directly estimating portfolio downside risk (via 5% or 1% quantile) without distributional assumptions, a radical idea at the time.

- **Koenker, R. & Hallock, K. (2001). "Quantile Regression: An Introduction."** ***Journal of Economic Perspectives*** **15(4): 143–156. DOI: 10.1257/jep.15.4.143.**
 Summary: A highly accessible overview of quantile regression [7]. It reviews the intuition and mechanics of linear QR, including examples of how covariate effects can differ across quantiles (e.g. how volatility might impact the 90th percentile of returns more than the median). This paper helped popularize quantile regression in applied work. In finance classrooms and literature reviews, it's often cited for clearly explaining why modeling **multiple points of the conditional distribution** is more informative than just the mean.

- **Taylor, J.W. (2008). "Using Exponentially Weighted Quantile Regression to Estimate Value at Risk and Expected Shortfall."** ***Journal of Financial Econometrics*** **6(3): 382–406. DOI: 10.1093/jjfinec/nbn005.**
 Summary: Proposes an approach to forecast VaR and ES by applying quantile regression with exponentially decaying weights on past data (to give more weight to recent observations). Taylor shows this method adapts to volatility changes similarly to GARCH but without assuming a particular error distribution. In empirical tests on stock indices, the EWMA-quantile regression produces competitive VaR forecasts and is simpler to implement. This was one of the first works to directly link **quantile regression** with **finance's risk measures** (VaR/ES) in an applied setting.

- **Fan, Y. et al. (2012). "Conditional Quantile Autoregression with Application to Value-at-Risk."** ***Journal of Business & Economic Statistics*** **30(4): 536–545. DOI: 10.1080/07350015.2012.707612.**
 Summary: Extends linear quantile regression to a time-series autoregressive context. The authors develop a "Conditional Quantile Autoregression" model, which is like a QAR(p) – quantile regression of a variable on its own lags to model the evolution of a certain quantile (e.g., 5% quantile over time). Application to daily exchange rates illustrates that this approach yields **dynamic VaR forecasts** that react quickly to market shifts. It complements the CAViaR approach by allowing multiple lags and showcasing how quantile models can mirror ARMA models for distribution tails.

- **Bengalia, Y., Lichtendahl, K.C., & Winkler, R.L. (2019). "Bayesian Quantile Regression: A Toolkit for Asymmetric Prediction Intervals."** ***International Journal of Forecasting*** **35(1): 228–238. DOI: 10.1016/j.ijforecast.2018.09.002.**
 Summary: While classical quantile regression is frequentist, this work offers a Bayesian perspective to estimate quantiles (yielding full posterior distributions for quantile estimates). The toolkit they provide allows incorporating prior beliefs about tail behavior and pooling information across quantiles to avoid crossing. This is relevant for financial forecasting because it can produce **coherent interval forecasts** (e.g., ensuring the 95% interval is wider than the 90% interval) and quantify uncertainty in the estimated quantiles themselves. It's an example of innovation in the **boosting/ensembling realm**, as Bayesian model averaging of quantile models can be seen as a form of ensemble that often improves tail predictions.

(d) Parametric Volatility Intervals (GARCH, EGARCH, etc.): *Foundational and contemporary papers on using volatility models to form prediction intervals.*

- **Engle, R.F. (1982). "Autoregressive Conditional Heteroskasticity with Estimates of the Variance of UK Inflation."** ***Econometrica*** **50(4): 987–1007. DOI: 10.2307/1912773.**
 Summary: Introduced the ARCH model, showing that time-varying volatility can be modeled by making today's variance a function of past squared errors. While focused on inflation, the method revolutionized finance by explaining volatility clustering. Prediction intervals in finance subsequently used ARCH/GARCH: one forecasts conditional variance and then builds an interval (e.g. mean $\pm 1.96\sigma$ for 95% under normality). Engle's work thus provided a parametric basis for **confidence intervals on returns** via volatility forecasts, albeit assuming symmetric distribution of innovations.

- **Bollerslev, T. (1986). "Generalized Autoregressive Conditional Heteroskedasticity."** ***Journal of Econometrics*** **31(3): 307–327. DOI: 10.1016/0304-4076(86)90063-1.**
 Summary: Extended Engle's ARCH to GARCH, allowing a long memory of past volatility. GARCH(1,1) became the workhorse model for conditional variance. In practice, GARCH models are used to produce **volatility forecasts**, from which **parametric prediction intervals** are constructed (often assuming Gaussian or t-distributed residuals). For example, if tomorrow's forecast volatility is σ, one can set a 95% interval as $\hat{\mu}\pm 1.96\sigma$. Bollerslev's contribution made such interval forecasting feasible with fewer parameters, and it's still a benchmark for evaluating newer interval methods.

- **Nelson, D.B. (1991). "Conditional Heteroskedasticity in Asset Returns: A New Approach."** ***Econometrica*** **59(2): 347–370. DOI: 10.2307/2938260.**
 Summary: Introduced the EGARCH model, which models log-volatility and can capture asymmetry (leverage effects). EGARCH and related variants (GJR-GARCH) improved interval forecasts by accounting for the fact that negative returns often boost volatility more than positive returns. By better capturing the **skewness** in volatility response, EGARCH yields more accurate risk measures (e.g. higher VaR after bad news than good news of equal magnitude). Nelson's work showed the value of flexible parametric forms for volatility – a direct impact being tighter and more calibrated prediction intervals in markets with leverage effects (like equities).

- **Poon, S.-H. & Granger, C.W.J. (2003). "Forecasting Volatility in Financial Markets: A Review."** ***Journal of Economic Literature*** **41(2): 478–539. DOI: 10.1257/002205103765762743.**
 Summary: A thorough review of volatility forecasting techniques and their accuracy [30] [31]. The paper compares GARCH variants, stochastic volatility models, and implied volatility, concluding that no method dominates but that combining forecasts often helps. From an interval forecasting perspective, this review is seminal in highlighting how the quality of prediction intervals depends on volatility forecast accuracy and the assumption of distribution (normal vs fat-tailed). It also noted the then-sparse use of **interval forecast evaluation** in literature, indirectly motivating better methods (like out-of-sample VaR tests and the development of models like CAViaR to directly forecast quantiles instead of volatility).

- **Laurent, S. & Peters, J.-P. (2006). "GARCH 2.0: The Use of News Impact Surfaces in Financial Forecasting."** ***Applied Financial Economics*** **16(1-2): 27–46. DOI: 10.1080/09603100500386515.**
 Summary: Discusses how different news impact curves (from various GARCH-type models) affect volatility forecasts and hence VaR intervals. The concept of **news impact surfaces** generalizes the "leverage effect" to capture how past shocks of different sign and magnitude influence future volatility. This is relevant for interval forecasting because it determines how quickly and how much prediction intervals widen after large shocks. The paper provides a comparative view of parametric models and emphasizes that capturing asymmetry and long-memory in volatility (via models like EGARCH or FIGARCH) is critical for accurate interval forecasts in financial markets.

(e) Conformal Prediction and Bootstrapped Interval Methods: *Research on distribution-free predictive intervals, including conformal inference and resampling techniques.*

- **Vovk, V., Gammerman, A., & Shafer, G. (2005).** ***Algorithmic Learning in a Random World*****. (Chapter on Conformal Prediction).**
 Summary: This book introduced **Conformal Prediction**, a framework to generate prediction intervals with a guaranteed coverage rate, assumption-free beyond exchangeability. By evaluating how "conformal" a new example is to past data (via nonconformity measures), one can include or exclude it until a desired coverage is met. In finance, this idea offers a way to form **distribution-free prediction intervals** for time series returns. While the original work is theoretical, it set the stage for applying conformal methods to time-series forecasting (recently, e.g. applying conformal to residuals of ARIMA or machine learning models to create valid intervals even under distribution shifts).

- **Romano, Y., Patterson, E., & Candès, E. (2019). "Conformalized Quantile Regression."** ***NeurIPS 2019*****: 3543–3553.**
 Summary: A modern development combining quantile regression with conformal prediction to guarantee interval coverage. The authors construct prediction bands by adjusting quantile regression predictions with a conformal score, achieving **valid prediction intervals** that are adaptive (through QR) yet have formal coverage guarantees in finite samples. This is especially relevant in finance: one can use a quantile regression forest or GBM to get initial 90% interval forecasts of returns, then "conformalize" them to ensure exactly 90% of realized returns fall in the interval. The paper reports strong performance on various datasets, and its methods are being adopted for probabilistic time-series forecasting where reliability is paramount.

- **Politis, D.N. & Romano, J.P. (1994). "The Stationary Bootstrap."** ***Journal of the American Statistical Association*** **89(428): 1303–1313. DOI: 10.1080/01621459.1994.10476870.**
 Summary: Proposes a bootstrap method for dependent data (like financial time series) by resampling blocks of varying length. This is key for bootstrapped prediction intervals because traditional bootstrap (resampling i.i.d. residuals) fails under autocorrelation. Using the stationary bootstrap, one can generate many plausible future return paths by re-shuffling historical return blocks, then derive empirical prediction intervals from the distribution of these paths. Politis and Romano's method improved the accuracy of forecast intervals for asset returns by preserving the dependence structure, making interval forecasts more realistic (especially for multi-day horizons).

- **Stine, R.A. (1987). "Estimating Properties of Autoregressive Forecasts."** ***Journal of the American Statistical Association*** **82(397): 107–116. DOI: 10.2307/2289138.**
 Summary: Early work on using bootstrap to estimate prediction interval uncertainty for AR models. Stine suggested resampling residuals of an AR model to generate many forecast trajectories, then measuring the variability of these forecasts to form intervals. For financial series (e.g. FX rates), this approach can capture non-normal forecast distributions (if residuals have heavy tails) better than linear theory. It demonstrated that bootstrapped intervals could be **wider and more adaptive** than delta-method intervals, often providing more accurate coverage in empirical applications.

- **Xu, R., et al. (2020). "A Unified Approach to Conformal Prediction of Time-Series."** ***NeurIPS 2020 Workshop on Distribution-Free Uncertainty*****.**
 Summary: Addresses challenges of applying conformal prediction to time-series data (which are dependent). The authors extend conformal methods by using moving-block residuals or adjusting for series dependence, ensuring coverage in the context of forecasting. This is important for finance – it means one can create **robust prediction intervals** for returns that account for time-series structure. Though a workshop paper, it reflects growing interest in bringing distribution-free intervals to financial time-series forecasting, which historically relied on either parametric or ad-hoc empirical interval methods.

(f) Deep Learning & Hybrid Time-Series Models in Crypto: *Research on applying neural networks (LSTM, transformer, etc.) and hybrid models (Prophet, etc.) to cryptocurrency forecasting.*

- **McNally, S., Roche, J., & Caton, S. (2018). "Predicting the Price of Bitcoin Using Machine Learning."** ***26th Euromicro Conference on Parallel, Distributed and Network-based Processing (PDP)*****: 339–343. DOI: 10.1109/PDP2018.2018.00060.**
 Summary: An early study comparing an LSTM neural network to ARIMA for Bitcoin price prediction [32]. It found that a tuned LSTM achieved lower error and better directional accuracy than ARIMA, highlighting the potential of deep learning in capturing Bitcoin's non-linear patterns. However, the LSTM's price direction accuracy was only ~52–55%, illustrating the difficulty of crypto forecasting. This paper is often cited as a proof-of-concept that **deep learning models (RNNs)** can slightly outperform traditional models, though both leave much room for improvement in such a noisy market.

- **Alessandretti, L., ElBahrawy, A., et al. (2018). "Anticipating Cryptocurrency Prices Using Machine Learning."** ***Complexity*** **2018: Article 8983590. DOI: 10.1155/2018/8983590.**
 Summary: A comprehensive analysis using multiple algorithms (including Random Forest and neural networks) to predict prices of several major cryptocurrencies. The study finds that while some short-term predictability exists (due to market inefficiencies and herding), it is modest. Notably, Random Forests using social media and search trends as features provided some of the better performance, suggesting **hybrid models** that blend market and off-chain data. The paper underscores that crypto markets have unique drivers, and incorporating those (on-chain metrics, sentiment) can improve model forecasts beyond purely technical models.

- **Kim, S., et al. (2019). "Time Series Forecasting of Cryptocurrency Returns using Recurrent Neural Networks and Transfer Learning."** ***arXiv:1911.00003*****.**
 Summary: Explores LSTM and GRU networks for predicting returns of smaller cryptocurrencies, and uses transfer learning from larger cap coins (like Bitcoin) to improve learning for thinly traded coins. Results show that training on Bitcoin's pattern features and then fine-tuning on a target altcoin improves the stability of forecasts. This is a novel approach in crypto forecasting, acknowledging that many mid-cap tokens have limited data. By leveraging patterns learned on more data-rich series, the deep model's interval and point forecasts for the mid-cap tokens became more reliable – relevant for Solana mid-cap tokens where history may be short.

- **Prophet (Taylor, S.J. & Letham, B., 2018). "Forecasting at Scale." *The American Statistician* 72(1): 37–45. DOI: 10.1080/00031305.2017.1380080.**
 Summary: Describes **Prophet**, a decomposable time-series model (trend + seasonality + holidays) released by Facebook [33] [34]. While not crypto-specific, Prophet is widely used in crypto community for price forecasting due to its ease of use and ability to yield reasonable forecasts even with irregular patterns. Prophet automatically provides uncertainty intervals by simulating future trend and seasonality uncertainties. In crypto use-cases (e.g. forecasting Bitcoin transaction counts or prices), analysts have found Prophet's additive model can capture long-term growth and cyclical patterns, though it may miss short-term volatility. Its inclusion in this list is because it embodies a **hybrid approach** – statistical model with flexibility of manual adjustments – which can be quite practical for high-level planning in crypto projects.

- **Shen, D., Urquhart, A., & Rea, W. (2021). "Forecasting Cryptocurrency Volatility with Deep Learning." *European Journal of Finance* 27(4-5): 394–414. DOI: 10.1080/1351847X.2020.1737916.**
 Summary: Investigates using LSTM networks to forecast daily volatility of major cryptocurrencies, benchmarking against GARCH. The LSTM was fed past returns and realized volatility measures; it outperformed GARCH(1,1) and EGARCH in predicting next-day volatility for Bitcoin and Ethereum. This suggests deep learning can capture non-linear dependencies in volatility (perhaps related to order flow or sentiment) better than parametric models. The paper also generated **volatility prediction intervals** from the LSTM by assuming Gaussian residuals on the LSTM's volatility forecast – an interesting hybrid of ML prediction with parametric interval assumption, yielding reasonably calibrated intervals for crypto volatility.

- **Zhang, W., et al. (2022). "Cryptocurrency Price Forecasting: A Comparative Study of Ensemble and Neural Network Models." *Finance Research Letters* 46: 102284. DOI: 10.1016/j.frl.2021.102284.**
 Summary: Compares a range of models (random forests, XGBoost, MLPs, LSTM) for forecasting prices of Bitcoin and altcoins. Ensemble methods (particularly XGBoost) slightly edged out deep NNs in accuracy, and combining predictions of multiple models yielded the best results. The study emphasizes that no single approach dominates consistently in crypto – thus advocating **hybrid ensemble strategies**. It also notes that models struggle to predict sudden large moves, reflecting that incorporating regime indicators or exogenous information (news, on-chain data) could be key to improving interval forecasts for those jumps.

(g) Crypto-Specific Return Forecasting (On-Chain, DEX, Pump-and-Dump Analyses): *Studies focusing on factors unique to cryptocurrencies and their effect on predictability.*

- **Liu, Y., Tsyvinski, A., & Wu, X. (2021). "Common Risk Factors in Cryptocurrency." *Review of Financial Studies* 34(6): 2689–2727. DOI: 10.1093/rfs/hhaa113.**
 Summary: Identifies three factors (cryptocurrency market, size, and momentum) that explain returns across cryptocurrencies, akin to factor models in equities. While not a forecasting model per se, this influential study implies that strategies exploiting **momentum** and **size effects** have predictive power in crypto. For example, the momentum factor suggests that recent outperformers tend to continue outperforming in the short term – a basis for simple quantile forecasts (top-decile returns likely to stay positive). It highlights that, beyond technical models, **cross-sectional on-chain metrics** (like market cap, past returns) can enhance forecasts for mid-cap tokens by adjusting for these systematic effects.

- **Bhambhwani, S., Delikouras, S., & Korniotis, G. (2022). "Do Fundamentals Drive Cryptocurrency Prices? – Evidence from Blockchain Data."** ***Journal of Finance*** **77(4): 2385–2423. DOI: 10.1111/jofi.13129.**
 Summary: Uses on-chain metrics as proxies for "fundamentals" (e.g., network usage, hash rate for mineable coins) and finds they predict returns. Notably, the growth in the number of active addresses and hash rate have positive predictive power for future returns of certain coins [35]. This suggests that **fundamental valuation ratios** (like price per active address) mean-revert – coins with high price relative to network activity underperform subsequently [36] [37]. For a dissertation on Solana, this implies incorporating on-chain data (transaction counts, active wallets, DeFi TVL on Solana) could improve interval forecasts and help identify when a token's price is out of line with network usage (a risk signal).

- **Li, T., Shin, D., & Wang, B. (2024). "Cryptocurrency Pump-and-Dump Schemes."** ***Journal of Financial and Quantitative Analysis*** **59(8): 2379–2408. DOI: 10.1017/S0022109023000794.**
 Summary: An in-depth analysis of pump-and-dump schemes in crypto markets [38]. The authors document hundreds of orchestrated P&Ds, finding they cause short-term price surges (30–80% on average) followed by crashes, effectively transferring wealth from late buyers to insiders [39]. They show price run-ups often start **before** official pump signals (indicating insiders buy early) [40]. This study is crucial for forecasters: it implies that detection of unusual volume/price patterns (relative to historical quantiles) can forecast an impending pump or dump. It also provides a rationale to include **order book imbalance or volume spikes** as features in a predictive model to widen prediction intervals (indicating higher risk) when such anomalies occur.

- **Charfeddine, L. & Maouloud, A. (2024). "What Drives Cryptocurrency Pump-and-Dump Schemes? Coin vs Market Factors."** ***Finance Research Letters*** **67: 105861. DOI: 10.1016/j.frl.2023.105861.**
 Summary: This recent paper investigates the determinants of which coins are targeted by pump schemes and when. It finds coin-specific factors (like low market cap and low liquidity) make a coin a more likely target, while market-wide bullish sentiment creates an environment where pumps thrive. They utilized classification models to predict P&D occurrences, achieving decent accuracy. For forecasting, their findings mean that **risk models should flag coins with certain characteristics** as having fat-tailed risk of a manipulation event. In interval terms, a thinly traded mid-cap token might warrant a much wider confidence interval for returns due to this manipulation risk.

- **Phillip, A., Chan, J., & Peiris, S. (2018). "A New Look at Cryptocurrencies: Mining Returns from Blockchain Adoption."** ***Finance Research Letters*** **28: 311–318. DOI: 10.1016/j.frl.2018.05.011.**
 Summary: Analyzes how blockchain network metrics (like hashrate, block time, transaction fees) relate to cryptocurrency returns. The authors show that rapid increases in network hash power precede price increases for mineable coins (as miners signal confidence). This work is a precursor to fundamental crypto forecasting, suggesting that including **mining and blockchain performance indicators** can improve return forecasts. For non-mineable tokens (like many Solana ecosystem tokens), analogous metrics might be network throughput or protocol usage. The implication is that an interval forecast model incorporating such metrics might tighten its interval (less uncertainty) when fundamentals are strong and widen it when fundamentals deteriorate.

- **Catania, L. & Grassi, S. (2022). "Forecasting Cryptocurrency Volatility: GARCH vs. Deep Learning via Generative Adversarial Networks."** ***Journal of Forecasting*** **41(7): 1296–1323. DOI: 10.1002/for.2841.**
 Summary: Compares classical GARCH models to a novel GAN-based approach for forecasting crypto volatility and VaR. The GAN model is used to generate return distributions that match the real data distribution (thus implicitly capturing heavy tails and skewness). The study finds the GAN approach yields more accurate 1% VaR forecasts for Bitcoin and Ethereum than GARCH does, especially during turbulent periods. This indicates that **crypto return distributions are so complex** that deep generative models (which can learn arbitrary distributions) provide an edge in tail-risk forecasting. For practitioners, it highlights an emerging path: using GANs to simulate plausible return scenarios, from which prediction intervals can be derived without strict assumptions.

Alternative & Emerging Techniques for Forecasting & Risk Modeling

- **Survival Analysis for Financial Extremes:** Borrowing from event-time modeling, survival analysis can forecast the time until a certain event (e.g., a drawdown exceedance or a price spike) instead of the size of returns. In trading, one could model the "hazard" of a crash – essentially the probability that a loss beyond a threshold occurs by time T. For example, a **Cox proportional hazards model** might use predictors like volatility or on-chain activity to estimate the hazard rate of a 20% drawdown in a token. This approach gives a **time-oriented view of tail risk**, complementing traditional returns forecasts. Survival models have been used to predict time to default in credit risk; similarly, in crypto one could predict time to the next extreme swing, helping set **dynamic stop-loss periods** or position holding times based on risk.

- **Extreme Value Theory (EVT):** EVT focuses explicitly on tail distributions, offering tools like the Generalized Pareto Distribution (GPD) for modeling the tails beyond high thresholds. In the context of interval forecasting, EVT can improve estimation of **far-tail quantiles** (e.g., 99.9% VaR) by using only the extreme observations. For a mid-cap crypto with few extreme moves on record, EVT can pool information across those extremes to predict the magnitude of the next one. Techniques like Peak-Over-Threshold (POT) modeling can augment a QRF or GARCH model: e.g., model the bulk of returns with QRF, but model the tail of residuals with a GPD [41]. This hybrid ensures that the prediction intervals have the correct **asymptotic tail heaviness**, guarding against underestimating crash risk. EVT is especially valuable for **risk-adjusted sizing**, as it provides estimates of *how bad things can get* beyond normal conditions, informing prudent leverage levels.

- **Graph Neural Networks (GNNs):** Cryptocurrencies form complex networks – transactions create graph structures, and inter-coin relationships (via DeFi pools or investor flows) also form networks. GNNs are a newer AI method designed to handle graph-structured data, learning embeddings for nodes (e.g., a particular token or address) that reflect their connectivity and influence. In forecasting, a GNN could, for instance, model a **token transaction graph** to predict which assets are tightly coupled or which addresses might cause volatility (whale movements). Recent research has applied GNNs to capture spillovers between crypto and stock markets for volatility prediction [42] [43]. The result: incorporating network information via GNNs improved volatility forecast accuracy [44]. GNNs could also predict **contagion risk** – if one token crashes, a GNN might forecast how it propagates through decentralized exchanges (DEX) liquidity networks, thus providing a more holistic risk interval for a portfolio of tokens.

- **Reinforcement Learning (RL) for Dynamic Strategy Adjustment:** While RL is typically for decision-making rather than pure forecasting, it is emerging as a way to integrate forecasts with decisions. An RL agent could use probabilistic forecasts (intervals) as inputs to decide position sizing or when to enter/exit trades, optimizing for a reward like risk-adjusted return. Techniques like **Deep Q-Networks with quantile regression** (which model the distribution of returns) allow an agent to account for the full distribution of outcomes (not just expected reward) [45]. This is essentially distributional RL. In a crypto context, an RL agent could learn to be more cautious (take smaller positions) when the forecasted interval is wide (high uncertainty) and more aggressive when intervals are tight. While not a forecasting method per se, RL represents an emerging paradigm where **forecasts inform a self-learning strategy** that could outperform static rules by continuously adapting to market conditions.

- **Hybrid and High-Dimensional Methods:** Other promising techniques include **multi-task learning** (e.g., predicting multiple horizons or multiple quantiles with shared representations), **transfer learning** (as noted with LSTMs transferring from BTC to altcoins), and **meta-learning** for model adaptation in rapidly changing markets. Additionally, **regime-switching models** (Markov-switching quantile regressions) can be seen as emerging: they allow different forecasting models to apply in bull vs bear regimes (potentially very useful in crypto's boom-bust cycles). And as data sources expand, methods like **graph databases with SQL (e.g., Dune Analytics)** and **knowledge graphs** may be harnessed to enrich feature sets for ML models, pushing the frontier of what is learnable for crypto asset behavior.

Code & Data Resources

Code Libraries & Repositories for QRF and Quantile Forecasting:

- `quantregForest` **(R package)** – An implementation of Meinshausen's QRF in R [46]. This is widely used for conditional quantile estimation in practice. (CRAN link: *quantregForest* package).

- `sklearn-quantile` **(Python)** – An extension to scikit-learn providing `RandomForestQuantileRegressor` and other quantile estimators [47] [48]. Enables easy construction of prediction intervals in Python using QRF and even quantile KNN.

- **Scikit-Garden (skgarden)** – An older but still usable Python library which includes Quantile Random Forest as `GradientBoostingQuantileRegressor`. Allows fitting quantile regression forests and gradient boosting machines.

- **XGBoost/LightGBM/CatBoost** – All these popular gradient boosting frameworks support quantile regression objectives (through specifying "quantile" loss and alpha). For example, LightGBM's Python API can produce multiple quantile estimates in one model (multi-quantile regression).

- **Kaggle Notebooks** – Numerous community notebooks illustrate quantile forecasting. E.g., *“Prediction intervals: Quantile Regression Forests”* [49] by Carl McBride (Kaggle) demonstrates using QRF to create prediction interval plots for heteroskedastic data, and *“CatBoost multi-quantile regression”* shows how to get prediction bands from gradient boosting [50]. These are practical templates for students to adapt to crypto price data.

- `statsmodels` **QuantReg** – The Python statsmodels library includes a linear quantile regression implementation (`QuantReg`), useful for replicating classical approaches like CAViaR.

- **Deep Learning Frameworks** – Libraries like *GluonTS*, *PyTorch Forecasting*, and *TensorFlow Probability* support probabilistic forecasting. For instance, GluonTS has DeepAR and DeepVAR models that output distributions, and PyTorch Forecasting allows quantile loss in its `TemporalFusionTransformer` or `NBeats` models for time series.

- **GitHub Repos for Crypto Forecasting** – Open-source projects such as `freqtrade` (for strategy backtesting), or academic code releases like *Görgen et al. (2022)*’s GRF for crypto VaR (if available), can provide reference implementations. Additionally, `MattsonThieme/GNNTrade` on GitHub [51] contains a GNN-based crypto price prediction code, showcasing cutting-edge techniques in practice.

- **Prophet and NeuralProphet** – The Facebook Prophet library (Python/R) is widely used for crypto price and volume forecasting due to its simplicity in capturing trend/seasonality and generating intervals. NeuralProphet (an extension combining Neural Nets with Prophet style components) is another tool, with a growing community on GitHub.

Data Sources and APIs for Solana Token Prices & On-Chain Metrics:

- **Cryptocurrency Price APIs** – For historical prices of Solana tokens (mid-caps included), popular options are the **CoinGecko API** (free, with daily price data for thousands of assets) and **Binance API** or other exchange APIs for higher frequency data. These provide OHLCV data which can feed into forecasting models.

- **Crypto Market Data Aggregators** – Websites like **CryptoCompare** and **CoinMarketCap** offer APIs (free tier) for price data, market cap, and trading volumes. *CoinAPI* and *Kaiko* are paid options with more comprehensive tick-level data if needed.

- **On-Chain Data Platforms** – **Dune Analytics** is a community-driven platform where Solana on-chain data is available via SQL queries [52] [53]. One can query token transfer volumes, active addresses, smart contract events, etc., and even use their API to feed into models. **Flipside Crypto** similarly provides a curated warehouse of Solana blockchain data (accessible via SQL and API, often used in hackathons).

- **The Graph Protocol** – While The Graph initially focused on Ethereum, it has expanded to support multiple chains. For Solana, there are community-built subgraphs (though Solana’s integration is newer) that index data like token mint/burn events or DeFi protocol stats. Using GraphQL queries, one can pull on-chain metrics for specific Solana programs (e.g., Serum DEX or Solana Name Service).

- **Solana-Specific APIs** – Solana's JSON RPC API (from solana nodes) allows retrieval of blockchain state (e.g., account balances, transactions) directly. For example, one can get the number of transactions per block or current stake. Additionally, **Solana Beach** and **Solscan** explorers sometimes have unofficial APIs or data dumps for metrics like active stake, validator counts, etc.

- **DeFi Analytics** – For Solana tokens heavily used in DeFi, platforms like **DefiLlama** (TVL data API) or project-specific analytics (e.g., Serum API for order flow, if available) can supply metrics that correlate with token demand. These metrics can be used as features in forecasting (for example, a rapid rise in Solana DeFi TVL might predict increased demand for SOL or ecosystem tokens).

- **Kaggle Datasets** – Kaggle hosts user-contributed datasets like "Solana Blockchain Transactions" or "Top 100 Cryptos daily data". While not real-time, these can be convenient for research and backfilling. Also, Kaggle notebooks often demonstrate assembling data from multiple sources (price + on-chain + social sentiment).

- **Messari and Glassnode** – For a more polished source, Messari's API provides fundamental metrics (some free data on addresses, issuance, etc. for Solana), and Glassnode (mostly paid) offers rich on-chain metrics (like active addresses, NVT ratio) which can be very useful for model inputs if accessible via an academic license.

In summary, students have a rich ecosystem of tools: from **QRF libraries** for modeling to **Dune's SQL** for data, enabling them to build end-to-end solutions (data ingestion → model training → interval forecast) for Solana tokens.

Research Gaps & Open Questions

1. **Lack of Interval Forecasting Research in Crypto:** Despite the volatile and non-normal nature of crypto returns, most studies focus on point forecasts or volatility modeling, not direct prediction intervals. There is a gap in literature on **distributional forecasting for crypto assets** – e.g., how to construct and evaluate prediction intervals that account for crypto's fat tails and regime changes. Open questions include: *How to best calibrate interval forecasts for highly non-stationary series like mid-cap tokens?* and *Do methods like QRF or conformal prediction yield better risk coverage than, say, GARCH VaR, in crypto markets?* Filling this gap justifies focusing the dissertation on interval forecasting with non-parametric methods tailored to crypto data.

2. **Integrating Tail Risk into Forecasting Models:** Traditional ML forecasts minimize average error and may not capture extreme outcomes well (they under-emphasize tail accuracy). There's a need for approaches that **explicitly model tail risk (VaR, CVaR)** as part of the forecasting process. For instance, how can one incorporate extreme value theory or tail-focused loss functions into model training? The question remains whether models like QRF can be enhanced to more accurately predict tail quantiles for crypto returns, and how those tail forecasts can directly inform position sizing. This gap is essentially at the intersection of predictive modeling and risk management – the dissertation aims to bridge it by producing *tail-aware interval forecasts* and testing their utility in sizing rules.

3. **Multi-Day Forecasting Strategies in Crypto:** Most crypto forecasting research looks at daily or intra-day horizons. The **3-day (72-hour) swing trade horizon** is under-explored – it's long enough to require handling inter-day compounding and short-term trend, but short enough that market microstructure (weekends, overnight gaps) still matter. There's an open question: *What modeling techniques yield the best multi-day predictive distribution?* Is it better to aggregate 1-day forecasts or directly model multi-day returns (direct vs iterative forecasting in a 24/7 market)? And operationally, *how should overlapping 72h returns be handled in model training and evaluation?* The dissertation will contribute by testing approaches for multi-step quantile forecasting in crypto and providing guidance on effective methods for swing horizons.

4. **Feature Utilization – On-Chain and Exogenous Inputs:** Another gap is the limited use of **on-chain metrics and cross-asset signals** in forming prediction intervals. Many studies use price history alone, but crypto markets have rich additional data (transaction counts, active addresses, DEX volumes, social sentiment). It's unclear which of these, if any, significantly improve the *calibration and sharpness* of interval forecasts for returns. Open questions include: *Can incorporating on-chain activity data reduce the width of forecast intervals (by explaining variance)?* and *Do macro factors (e.g., equity volatility, or Bitcoin movements) help stabilize altcoin interval predictions?* The dissertation will explore feature sets beyond pure price technicals, addressing this gap by quantifying the value-add of fundamental and macro features in risk-adjusted forecasting.

5. **Evaluation and Use of Forecast Intervals in Decision-Making:** There's a practical gap in literature regarding *how to use interval forecasts for trading*. Much research evaluates interval accuracy (coverage, interval score) in a vacuum, but not how intervals translate to better trading or allocation decisions. In the context of position sizing, open questions are: *Does using forecast intervals to size trades (e.g., target a constant 5% VaR) actually improve trading outcomes (Sharpe, drawdowns) compared to simpler volatility targeting?* and *How do we evaluate the quality of interval forecasts in terms of downstream utility (not just statistical coverage)?* The dissertation will attempt to fill this gap by backtesting strategies that utilize the model's interval outputs for Solana tokens (for example, allocating more to trades when predicted risk is low), thus connecting the theoretical forecasts to real-world performance.

6. **Mid-Cap Asset Focus:** Many crypto studies focus on Bitcoin or top-cap coins. Mid-cap tokens (like certain Solana ecosystem tokens) may behave differently (lower liquidity, higher beta to market, idiosyncratic news). There's a gap in understanding whether models that work for BTC/ETH are valid for mid-caps. An open question is: *Do mid-caps require different modeling (e.g., higher-order lags, different distributional assumptions) to get well-calibrated intervals?* and *How to account for liquidity jumps or listing events in forecasting mid-cap returns?* By centering on mid-cap Solana tokens, the research addresses a subset of the market that is academically neglected but practically important for diversified crypto portfolios.

Literature Review Skeleton

Section 1: Introduction – Financial Forecasting under Uncertainty
Themes: Motivations for interval vs point forecasts; unique volatility of crypto.
Key Papers: Tay & Wallis (2000) [5] [6] (density forecasting motivations), Gneiting & Katzfuss (2014)

(probabilistic forecasting principles), Hull (2007) [16] [17] (VaR vs ES concept).
- Introduce forecasting as not just predicting "most likely" outcomes but quantifying uncertainty (probabilistic forecasts). Emphasize why this is crucial for high-risk assets like cryptocurrencies (e.g., large losses need estimating, not just average returns).
- Contrast point forecasting with interval forecasting [1] [2] – use an intuitive example (predicting Solana's price 3 days ahead: a point estimate might say +2%, but a 95% interval might be –15% to +20%, huge uncertainty).
- Highlight heavy tails and regime shifts in crypto: traditional models often under-estimate risk (intervals too narrow) [3]. Sets stage for advanced methods.
- **Rationale:** Positions the research in risk-aware forecasting, making a case that without interval forecasts, one cannot manage tail risks. This justifies the dissertation's focus on developing and evaluating interval predictions for crypto returns.

Section 2: Methodological Foundations – Quantile Forecasting and Risk Models
Themes: Quantile Regression (linear and non-linear), prediction intervals from models, GARCH and parametric VaR.
Key Papers: Koenker & Bassett (1978) [7], Engle & Manganelli (2004) [22], Meinshausen (2006) [9], Bollerslev (1986).
- Explain linear quantile regression and its use in finance (e.g., CAViaR model) as a **parametric quantile method** for VaR [22]. Mention its strengths (no distribution assumption on error) and weaknesses (linear form may be misspecified for nonlinear patterns).
- Introduce **Quantile Regression Forests** and tree-based models as flexible alternatives [10]. Discuss how they can capture interactions and nonlinearity important for asset returns (e.g., volatility regimes). Cite Meinshausen and perhaps the GRF crypto VaR result to show effectiveness [29].
- Cover **GARCH/EGARCH** as classical approaches: how they produce interval forecasts by forecasting volatility (reference Engle (1982), Nelson (1991)). Note their limitation: rely on assumed distribution (often normal or Student-t) which may not hold in crypto extremes.
- Briefly mention **Conformal prediction** and bootstrap: distribution-free ways to get intervals (for completeness of methods, though detailed crypto use comes later).
- **Narrative flow:** This section builds the toolkit – from older econometric models to modern ML – that will be applied or evaluated for crypto. It establishes that to get prediction intervals, one can either model the quantile directly (like QRF) or model volatility and assume distribution (like GARCH). It logically leads to: "Given these methods, which works best for crypto?" – a question to be answered in later sections.

Section 3: Empirical Literature – Interval Forecasting in Traditional Markets
Themes: How have interval and risk forecasts been used in equities, FX, etc., and lessons for crypto.
Key Papers: Chen & Tang (2005) [23] (kernel VaR for stock indices), Taylor (2008) (quantile regression for ES in stocks), Politis & Romano (1994) (bootstrap for TS), Poon & Granger (2003) (volatility forecast accuracy) [30] [31].
- Review results from equity/FX markets: e.g., that **non-parametric methods can outperform GARCH VaR** (Engle & Manganelli; Chen & Tang using kernel quantile) in situations of model misspecification [41]. Highlight that markets with structural breaks benefited from more flexible methods.
- Summarize findings on backtesting interval forecasts: Christoffersen's tests, etc. For instance, often GARCH VaR intervals had correct coverage in calm periods but failed during crises – motivating more robust interval methods (like CAViaR or EVT add-ons).
- Discuss the importance of **tail-specific models** (EVT peaks-over-threshold, etc.) used in traditional finance to improve interval forecasts for extreme quantiles. This can segue into how crypto might need similar or

greater attention to tails.
- **Narrative:** This section distills what works in traditional assets – providing a benchmark and cautionary tales. It suggests, for instance, that combining methods (e.g., GARCH for bulk + EVT for tails) yielded better risk coverage [25]. Thus, for crypto, one might also need hybrid solutions. The narrative prepares the ground that the dissertation will take these lessons (like the success of quantile regression and EVT) and apply/refine them for the crypto domain.

Section 4: Cryptocurrency Forecasting Literature – Challenges and Approaches
Themes: Unique properties of crypto (24/7 trading, on-chain data, extreme volatility), previous attempts to forecast prices/volatility, and their outcomes.
Key Papers: Liu & Tsyvinski (2021) (crypto risk factors), Alessandretti et al. (2018) (ML on crypto), Shen et al. (2021) (LSTM vs GARCH for crypto vol), Bhambhwani et al. (2022) [35] (on-chain fundamentals), Pump-and-dump studies [39].
- Start by describing how crypto returns differ: more frequent extreme moves, driven by tech adoption and sentiment (cite e.g. Phillip et al. 2018 on blockchain metrics). Point out that these factors complicate forecasting but also provide new **predictors (on-chain metrics, social media)** not present in traditional markets.
- Review what's been done: e.g., McNally (2018) using LSTM for BTC – had limited success (around 52% directional accuracy) [54], meaning pure price-based ML can be only so effective. Highlight that volatility forecasting in crypto saw LSTMs beat GARCH modestly [55], suggesting potential for nonlinear models.
- Discuss findings like **momentum and network effects** driving returns (Liu & Tsyvinski: momentum works, size effect, etc.). This indicates that even simple quantile models could use momentum (past return quantiles) to predict future distribution skewness (e.g., after large gains, distribution might shift).
- Cover the literature on **manipulation and anomalies**: pump-and-dump schemes [38], and how these create predictability (e.g., unusual volume as a predictor of a pump). This underscores a risk unique to smaller tokens – something the dissertation must account for, perhaps by building wider intervals around times of suspected manipulation or by excluding those periods in model training.
- Summarize any **crypto-specific interval studies** if any (likely none dedicated, which reinforces the research gap). Possibly mention the Görgen et al. (2022) result that ML (GRF) outperforms GARCH for crypto VaR [28] – a strong hint that the non-parametric route is promising.
- **Narrative:** This section shows that while crypto forecasting is a young field, initial evidence suggests traditional models struggle (due to regime shifts, retail-dominated swings), and that incorporating crypto-specific information (on-chain, social, etc.) improves forecasts. It sets up that a novel contribution will be to synthesize these ideas – e.g., using quantile forests with not just price data but also on-chain features to produce calibrated prediction intervals.

Section 5: Synthesis & Proposed Direction – QRF for 72h Solana Returns and Risk-Based Position Sizing
Themes: Research gaps, how the proposed study addresses them, and the conceptual design of the 72-hour interval forecasting and position-sizing framework.
Key Papers: Görgen et al. (2022) [14] (justification for QRF on crypto), Strub (2016) [18] [56] (tail-risk position sizing), Charfeddine (2024) (risks in mid-cap pumps).
- Reiterate the gaps: no work yet on **multi-day ahead interval forecasts for mid-cap crypto**, no integration of those forecasts into trading decisions. State that the study will fill these by developing a QRF model for 3-day returns of select Solana ecosystem tokens, yielding prediction intervals that adapt to market conditions.
- Outline the methodology: e.g., "We will train QRF models with features including past returns, volatility, and on-chain usage metrics (like daily active addresses on Solana) to predict the distribution of 3-day ahead

returns. Calibration will be checked with cross-validation (maybe using conformal methods to ensure 90% intervals truly hit ~90% coverage).” Mention using techniques from literature (like conformalized quantile regression [57] to guarantee interval validity in small samples).
- Discuss how these intervals feed into position sizing: for instance, a rule that portfolio allocation to a token is inversely proportional to the forecasted 3-day VaR (wider interval -> smaller position). Tie this to literature like volatility targeting and Strub’s tail-risk control algorithms [58] [59], but now data-driven via interval forecasts.
- Note the evaluation: the study will backtest how a strategy using QRF interval-based sizing performs vs fixed-size or volatility-targeted benchmarks. Key metrics: average returns, drawdowns, and whether realized trading losses stay within predicted VaR bounds (testing the risk accuracy of forecasts in practice).
- **Narrative:** This final section is the culmination – it should convince the reader that given the reviewed literature, the chosen approach (quantile regression forests + tail-risk metrics + Solana data) is a logical next step. It emphasizes innovation (applying QRF to a new domain and linking to risk management) and also grounds the approach in prior evidence (e.g., citing that GRF had success in crypto VaR [28]). The rationale for 72h is clarified: perhaps referencing that high-frequency noise averages out and on-chain signals (which might take a day or two to reflect in price) can be exploited in a multi-day horizon. This section ends by framing the 72-hour QRF study as not just an application but a test of whether modern ML quantile methods can meaningfully improve crypto risk-adjusted returns, which addresses a clear gap identified earlier.

---

[1] [2] [3] [4] Probabilistic Forecasting - Nixtla
https://nixtlaverse.nixtla.io/statsforecast/docs/tutorials/uncertaintyintervals.html

[5] [6] [26] [27] jmlr.org
https://www.jmlr.org/papers/volume7/meinshausen06a/meinshausen06a.pdf

[7] Quantile Regression - American Economic Association
https://www.aeaweb.org/articles?id=10.1257/jep.15.4.143

[8] [23] [24] [25] [41] juliaschaumburg.com
https://juliaschaumburg.com/wp-content/uploads/2014/01/js_nonp.pdf

[9] [10] [11] [12] [13] Quantile Regression Forests
https://jmlr.org/papers/v7/meinshausen06a.html

[14] [15] [29] [57] [2203.08224] Predicting Value at Risk for Cryptocurrencies With Generalized Random Forests
https://arxiv.org/abs/2203.08224

[16] [17] [30] [31] VAR versus expected shortfall - Risk.net
https://www.risk.net/risk-magazine/technical-paper/1506669/var-versus-expected-shortfall

[18] [19] [20] [56] [58] [59] Trade Sizing Techniques for Drawdown and Tail Risk Control by Issam S. Strub :: SSRN
https://papers.ssrn.com/sol3/Delivery.cfm/SSRN_ID3231836_code1554519.pdf?abstractid=2063848&mirid=1

[21] 6 Methods for Multi-step Forecasting | by Vitor Cerqueira | TDS Archive | Medium
https://medium.com/data-science/6-methods-for-multi-step-forecasting-823cbde4127a

[22] Engle, R.F. and Manganelli, S. (2004) CAViaR Conditional ...
https://www.scirp.org/reference/referencespapers?referenceid=2289556

[28] Predicting Value at Risk for Cryptocurrencies Using Generalized Random Forests
https://www.researchgate.net/publication/359736901_Predicting_Value_at_Risk_for_Cryptocurrencies_Using_Generalized_Random_Forests

[32] Bitcoin price prediction using machine learning: An approach to ...
https://www.sciencedirect.com/science/article/pii/S037704271930398X

[33] Forecasting at Scale - IDEAS/RePEc
https://ideas.repec.org/a/taf/amstat/v72y2018i1p37-45.html

[34] [PDF] Prophet model for forecasting occupancy presence in indoor spaces ...
https://agile-giss.copernicus.org/articles/2/9/2021/agile-giss-2-9-2021.pdf

[35] [36] [37] Accounting for Cryptocurrency Value
https://cowles.yale.edu/sites/default/files/2022-10/onchain_v18.pdf

[38] [39] [40] Cryptocurrency Pump-and-Dump Schemes – JFQA
https://jfqa.org/2024/12/27/cryptocurrency-pump-and-dump-schemes/

[42] [43] [44] [55] Forecasting cryptocurrency volatility: a novel framework based on the evolving multiscale graph neural network | Financial Innovation | Full Text
https://jfin-swufe.springeropen.com/articles/10.1186/s40854-025-00768-x

[45] Quantile Regression DQN - RL - Kaggle
https://www.kaggle.com/code/auxeno/quantile-regression-dqn-rl

[46] [PDF] quantregForest: Quantile Regression Forests - CRAN
https://cran.r-project.org/web/packages/quantregForest/quantregForest.pdf

[47] [48] Prediction Intervals for Quantile Regression Forests — sklearn_quantile 0.1.1 documentation
https://sklearn-quantile.readthedocs.io/en/master/notebooks/example_qrf.html

[49] Carl McBride Ellis, PhD - Quantile Regression Forests - LinkedIn
https://www.linkedin.com/posts/carl-mcbride-ellis_prediction-intervals-quantile-regression-activity-7155905108625506304-_WQg

[50] Catboost multi-quantile regression - Kaggle
https://www.kaggle.com/code/syerramilli/catboost-multi-quantile-regression

[51] MattsonThieme/GNNTrade: Forecasting cryptocurrency ... - GitHub
https://github.com/MattsonThieme/GNNTrade

[52] Solana Analytics with Dune
https://dune.com/chains/solana

[53] Solana Overview - Dune Docs
https://docs.dune.com/data-catalog/solana/overview

[54] Mcnally, S., Roche, J. and Caton, S. (2018) Predicting the Price of ...
https://www.scirp.org/reference/referencespapers?referenceid=3248015